

ADM-Assignment 3 -MD

Dhanush Myneni

2024-05-05

**

QA1. What is the difference between SVM with hard margin and soft margin?

Ans. Support Vector Machine (SVM) is a powerful supervised learning set of rules used for type and regression duties. The important distinction between SVM with difficult margin and soft margin lies in how they take care of the presence of outliers or noisy records factors.

Hard Margin SVM: Hard margin SVM goals to discover the maximum-margin hyperplane that separates the training perfectly without allowing any misclassifications. It assumes that the facts is flawlessly separable without any errors or outliers. If the facts isnt always linearly separable, hard margin SVM might not be able to discover a solution, and it will fail. Hard margin SVM is quite touchy to outliers because even a single outlier can save you the algorithm from finding a possible hyperplane.

Soft Margin SVM: Soft margin SVM is an extension of tough margin SVM that lets in for a few misclassifications or errors. It introduces a penalty parameter C . C that controls the exchange-off between maximizing the margin and minimizing the category blunders. Soft margin SVM can cope with non-linearly separable facts by way of allowing some information factors to be on the incorrect facet of the margin or even the wrong aspect of the hyperplane, for that reason making the set of rules greater sturdy to noisy information or outliers. The parameter C controls the penalty for misclassification. A smaller C cost allows for a bigger margin and extra misclassifications, whilst a larger C price reduces the margin to make sure fewer misclassifications. Soft margin SVM plays properly even if the information isnt always flawlessly separable or includes outliers, as it could find a compromise answer by adjusting the margin and permitting a few errors.

**

QA2.What is the role of the cost parameter, C , in SVM (with soft margin) classifiers?

Ans. In SVM with a soft margin, the cost parameter C plays an important role in striking a balance between the following - It affects the margin that the degree of complexity of choice boundary and the importance of error penalties are assessed. While a narrow C makes for a higher margin and more margin violations, it is less complicated and generalized. It does not necessarily attempt to become exact. In summary, a larger C penalizes strongly against misclassifications that are associated with less margin and a tighter boundary. Ultimately, the severity parameter should be compatible to the training data. As a result, C can be used to change the bias-variance trade-off appropriately chosen small parameters realize a smooth decision boundary while large parameters lend a more

complex boundary. Herein by adjusting C, the detrimental effects of outliers and noise are curbed, furthermore, guiding the working of the model among unlabeled data.

**

QA3. Will the following perceptron be activated (2.8 is the activation threshold)?

Ans. The activation function for the given inputs and weights is calculated as follows:

$$\text{Activation} = (0.1 \times 0.8) + (11.1 \times 0.2) = 0.08 - 2.22 = -2.14$$

The activation value is -2.14 which is lesser than the threshold value of 2.8 and the perceptron will not be activated in this scenario. Usually, perceptrons are activated if they are above the certain threshold and in this instance, the calculated value is lesser than the threshold value.

**

QA4. What is the role of alpha, the learning rate in the delta rule?

Ans. The delta rule is an widely used algorithm for being adjusted to the weights of a neural network during training. It calculates the difference between the predicted output and actual output which gives a hint of weights to be updated as well as the network performance improvement. The learning rate (α) is the most important among hyperparameters in the delta rule as it determines the steps size. A large learning rate brings about larger weight modifications, which will affect the model faster thereby allowing it to converge quicke through the ignorance of the right weights and the potential misbehavior of the algorithm, the risk of not achieving the optimal result may be present.

Conversely, a lower learning rate causes smaller updates and that makes the model to converge slowly but helps the model to generalize on unseen test datasets as well. Determining a right learning rate becomes essential for obtaining unparalleled accuracy. In Common one will start with a small value (e.g., 0.1 or 0.01) and then try different values to identify the optimum value. A balancing act is achieved by this process of iteration in such a manner that the network converges fast while at the same time not overconverge or coming to a rest at less than optimal solutions. Experimentation and tuning the learning rate are the prerequisites to getting the desired learning curve while building an optimal neural network.

**

Part B

#QB1. Build a linear SVM regression model to predict Sales based on all other attributes (Price, Advertising, Population, Age, Income and Education). Hint: use caret train() with method set to "svmLinear". What is the R-squared of the model?

```
library(ISLR)
library(dplyr)
```

```

## Warning: package 'dplyr' was built under R version 4.3.2
##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##     filter, lag
## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union

library(glmnet)

## Warning: package 'glmnet' was built under R version 4.3.3
## Loading required package: Matrix
## Warning: package 'Matrix' was built under R version 4.3.2
## Loaded glmnet 4.1-8

library(caret)

## Loading required package: ggplot2
## Warning: package 'ggplot2' was built under R version 4.3.2
## Loading required package: lattice

Filtered_Data <- Carseats %>% select("Sales", "Price",
  "Advertising", "Population", "Age", "Income", "Education")

set.seed(123)
train_index <- createDataPartition(Filtered_Data$Sales, p = 0.7, list =
  FALSE)
Training_Data <- Filtered_Data[train_index, ]
Testing_Data <- Filtered_Data[-train_index, ]

#Setting up the model.
SVM_Model <- train(Sales ~.,
  data = Training_Data,
  method = "svmLinear",
  trControl = trainControl(method = "cv", number
    = 10))

#Depicting Model's Summary

summary(SVM_Model)

## Length Class Mode
##      1  ksvm   S4

```

#Make predictions using the test dataset.

```
Predictions <- predict(SVM_Model, newdata = Testing_Data)
```

#Determining the R-squared value.

```
R_squared<- postResample(Predictions, Testing_Data$Sales)
```

```
R_squared
```

```
##      RMSE Rsquared      MAE
```

```
## 2.297064 0.385761 1.876610
```

#The R squared value of the model is .385

#QB2. Customize the search grid by checking the model's performance for C parameter of 0.1,.5,1 and 10 using 2 repeats of 5-fold cross validation.

```
library(caret)
```

```
Grid <- expand.grid(C = c(0.1,0.5,1,10))
```

```
Train_ctrl <- trainControl(method = "repeatedcv", number = 5, repeats = 2)
```

```
SVM_Linear_Grid <- train(Sales~., data = Filtered_Data, method = "svmLinear",  
                        trControl=Train_ctrl,  
                        preProcess = c("center", "scale"),  
                        tuneGrid = Grid,  
                        tuneLength = 10)
```

```
SVM_Linear_Grid
```

```
## Support Vector Machines with Linear Kernel
```

```
##
```

```
## 400 samples
```

```
## 6 predictor
```

```
##
```

```
## Pre-processing: centered (6), scaled (6)
```

```
## Resampling: Cross-Validated (5 fold, repeated 2 times)
```

```
## Summary of sample sizes: 319, 320, 320, 321, 320, 320, ...
```

```
## Resampling results across tuning parameters:
```

```
##
```

```
##      C      RMSE      Rsquared      MAE
```

```
##    0.1  2.270125  0.3649327  1.818221
```

```
##    0.5  2.269631  0.3642787  1.816916
```

```
##    1.0  2.269468  0.3642989  1.816430
```

```
##   10.0  2.269407  0.3643307  1.816397
```

```
##
```

```
## RMSE was used to select the optimal model using the smallest value.
```

```
## The final value used for the model was C = 10.
```

#QB3Train a neural network model to predict Sales based on all other attributes ("Price", "Advertising", "Population", "Age", "Income" and "Education")Hint: use caret train() with method set to "nnet". What is the R-square of the model with the best hyper parameters (using default caret search grid) - hint: don't forget to scale the data

#Data scaling

```
Scaled_data <- preProcess(Filtered_Data[-1], method = "scale")
Training <- predict(Scaled_data, Filtered_Data[-1])
```

```
#Train control
```

```
set.seed(27)
folds <- trainControl(method = "repeatedcv",
                      number = 5,
                      repeats = 2,
                      verboseIter = FALSE)
```

```
# Using default search grid, training the neural network model
```

```
set.seed(123)
NeuralNet_Cars <- train(Sales~., data = Filtered_Data ,
                      method = "nnet",
                      trControl = folds)
```

```
## # weights:  9
## initial  value 17583.375508
## final  value 15744.713000
## converged
## # weights:  25
## initial  value 18250.380996
## final  value 15744.713000
## converged
## # weights:  41
## initial  value 17969.872043
## final  value 15744.713000
## converged
## # weights:  9
## initial  value 17485.862634
## iter  10 value 15752.495629
## iter  20 value 15749.181183
## final  value 15749.165782
## converged
## # weights:  25
## initial  value 16731.281643
## iter  10 value 15750.726329
## iter  20 value 15747.342759
## iter  30 value 15747.245743
## iter  30 value 15747.245604
## iter  30 value 15747.245554
## final  value 15747.245554
## converged
## # weights:  41
## initial  value 17717.818624
## iter  10 value 15761.704139
## iter  20 value 15746.593674
## final  value 15746.528022
## converged
## # weights:  9
```

```
## initial value 18907.374924
## iter 10 value 15751.310870
## iter 20 value 15744.789068
## iter 20 value 15744.789038
## final value 15744.789038
## converged
## # weights: 25
## initial value 17691.631949
## iter 10 value 15766.322138
## iter 20 value 15744.962136
## iter 30 value 15744.727758
## iter 30 value 15744.727752
## iter 30 value 15744.727746
## final value 15744.727746
## converged
## # weights: 41
## initial value 18730.502990
## iter 10 value 15752.934200
## iter 20 value 15744.986373
## final value 15744.735157
## converged
## # weights: 9
## initial value 18150.262633
## final value 16133.992900
## converged
## # weights: 25
## initial value 18057.338235
## final value 16133.992900
## converged
## # weights: 41
## initial value 17359.482931
## final value 16133.992900
## converged
## # weights: 9
## initial value 17989.498119
## iter 10 value 16200.459178
## iter 20 value 16138.492102
## final value 16138.454408
## converged
## # weights: 25
## initial value 18745.698886
## iter 10 value 16153.006756
## iter 20 value 16137.256703
## iter 30 value 16136.548298
## final value 16136.530460
## converged
## # weights: 41
## initial value 17459.477458
## iter 10 value 16143.635989
## iter 20 value 16136.062684
```

```
## iter 30 value 16135.826166
## iter 40 value 16135.818629
## iter 40 value 16135.818513
## final value 16135.811494
## converged
## # weights: 9
## initial value 17665.772198
## iter 10 value 16142.809710
## iter 20 value 16134.094551
## iter 20 value 16134.094510
## final value 16134.094510
## converged
## # weights: 25
## initial value 17838.240242
## iter 10 value 16143.041882
## iter 20 value 16134.097228
## iter 30 value 16134.005971
## iter 30 value 16134.005867
## iter 30 value 16134.005862
## final value 16134.005862
## converged
## # weights: 41
## initial value 17513.611852
## iter 10 value 16144.718662
## iter 20 value 16134.116560
## final value 16134.008897
## converged
## # weights: 9
## initial value 18527.696365
## final value 16054.878300
## converged
## # weights: 25
## initial value 17788.524754
## final value 16054.878300
## converged
## # weights: 41
## initial value 17446.488348
## final value 16054.878300
## converged
## # weights: 9
## initial value 18790.788824
## iter 10 value 16060.229113
## iter 20 value 16059.341696
## final value 16059.340150
## converged
## # weights: 25
## initial value 18626.938869
## iter 10 value 16082.809489
## iter 20 value 16057.679412
## iter 30 value 16057.431616
```

```
## final value 16057.416615
## converged
## # weights: 41
## initial value 18338.133205
## iter 10 value 16153.067658
## iter 20 value 16058.367215
## iter 30 value 16057.518420
## iter 40 value 16056.957883
## iter 50 value 16056.696790
## iter 50 value 16056.696704
## iter 50 value 16056.696682
## final value 16056.696682
## converged
## # weights: 9
## initial value 17590.373920
## iter 10 value 16059.385823
## iter 20 value 16054.930268
## iter 20 value 16054.930247
## final value 16054.930247
## converged
## # weights: 25
## initial value 17195.152150
## iter 10 value 16063.994339
## iter 20 value 16054.983401
## final value 16054.887538
## converged
## # weights: 41
## initial value 19381.816748
## iter 10 value 16074.408053
## iter 20 value 16055.103463
## iter 30 value 16054.916332
## final value 16054.893441
## converged
## # weights: 9
## initial value 18955.393490
## final value 16058.233500
## converged
## # weights: 25
## initial value 18731.157319
## final value 16058.233500
## converged
## # weights: 41
## initial value 17519.676097
## final value 16058.233500
## converged
## # weights: 9
## initial value 18443.945205
## iter 10 value 16062.703228
## final value 16062.691766
## converged
```



```
## # weights: 25
## initial value 17324.504702
## iter 10 value 16063.061711
## iter 20 value 16060.975591
## iter 30 value 16060.769179
## iter 30 value 16060.769058
## iter 30 value 16060.768932
## final value 16060.768932
## converged
## # weights: 41
## initial value 17182.149602
## iter 10 value 16062.517552
## iter 20 value 16060.350351
## iter 30 value 16060.151610
## iter 40 value 16060.063571
## final value 16060.050628
## converged
## # weights: 9
## initial value 17573.438755
## iter 10 value 16062.663964
## iter 20 value 16058.284580
## iter 20 value 16058.284559
## final value 16058.284559
## converged
## # weights: 25
## initial value 17709.202204
## iter 10 value 16070.257094
## iter 20 value 16058.372123
## final value 16058.252642
## converged
## # weights: 41
## initial value 18558.603853
## iter 10 value 16075.463049
## iter 20 value 16058.432143
## final value 16058.249622
## converged
## # weights: 9
## initial value 18722.292207
## final value 16260.862700
## converged
## # weights: 25
## initial value 17799.426161
## final value 16260.862700
## converged
## # weights: 41
## initial value 18255.742928
## final value 16260.862700
## converged
## # weights: 9
## initial value 17481.607340
```

```
## iter 10 value 16273.839719
## iter 20 value 16265.331018
## final value 16265.323262
## converged
## # weights: 25
## initial value 18295.365403
## iter 10 value 16326.205795
## iter 20 value 16265.170871
## iter 30 value 16263.576040
## final value 16263.399405
## converged
## # weights: 41
## initial value 17990.022132
## iter 10 value 16263.841721
## iter 20 value 16262.695365
## final value 16262.680613
## converged
## # weights: 9
## initial value 18295.931560
## iter 10 value 16271.756840
## iter 20 value 16260.988301
## iter 20 value 16260.988251
## final value 16260.988251
## converged
## # weights: 25
## initial value 18603.949520
## iter 10 value 16276.382456
## iter 20 value 16261.041631
## iter 30 value 16260.871843
## iter 30 value 16260.871810
## final value 16260.870307
## converged
## # weights: 41
## initial value 18012.253619
## iter 10 value 16277.963842
## iter 20 value 16261.059863
## iter 30 value 16260.871990
## iter 30 value 16260.871973
## final value 16260.871512
## converged
## # weights: 9
## initial value 17699.831901
## final value 16275.959300
## converged
## # weights: 25
## initial value 19641.795202
## final value 16275.959300
## converged
## # weights: 41
## initial value 18752.645654
```

```
## final value 16275.959300
## converged
## # weights: 9
## initial value 17240.269435
## iter 10 value 16283.782698
## final value 16280.421142
## converged
## # weights: 25
## initial value 18553.272473
## iter 10 value 16279.424338
## iter 20 value 16278.558130
## iter 30 value 16278.527772
## final value 16278.496768
## converged
## # weights: 41
## initial value 17756.873961
## iter 10 value 16280.244531
## iter 20 value 16277.829378
## iter 30 value 16277.780248
## final value 16277.777787
## converged
## # weights: 9
## initial value 19121.132256
## iter 10 value 16284.843993
## iter 20 value 16276.061734
## iter 20 value 16276.061693
## final value 16276.061693
## converged
## # weights: 25
## initial value 18036.303840
## iter 10 value 16281.085424
## iter 20 value 16276.018400
## final value 16275.980929
## converged
## # weights: 41
## initial value 19355.400213
## iter 10 value 16288.718731
## iter 20 value 16276.106406
## final value 16275.970733
## converged
## # weights: 9
## initial value 18050.702295
## final value 15928.817500
## converged
## # weights: 25
## initial value 18673.513835
## final value 15928.817500
## converged
## # weights: 41
## initial value 18478.275519
```

```
## final value 15928.817500
## converged
## # weights: 9
## initial value 18032.649747
## iter 10 value 15935.432118
## final value 15933.274268
## converged
## # weights: 25
## initial value 18953.220557
## iter 10 value 15937.188701
## iter 20 value 15932.025175
## iter 30 value 15931.704036
## iter 40 value 15931.355821
## final value 15931.352118
## converged
## # weights: 41
## initial value 17775.320098
## iter 10 value 15931.261131
## iter 20 value 15930.730135
## iter 30 value 15930.634420
## iter 30 value 15930.634354
## iter 30 value 15930.634354
## final value 15930.634354
## converged
## # weights: 9
## initial value 17439.224458
## iter 10 value 15933.303079
## iter 20 value 15928.869215
## iter 20 value 15928.869195
## final value 15928.869195
## converged
## # weights: 25
## initial value 17831.381529
## iter 10 value 15939.629190
## iter 20 value 15928.942150
## final value 15928.827000
## converged
## # weights: 41
## initial value 16885.798107
## iter 10 value 15943.011273
## iter 20 value 15928.981143
## final value 15928.827662
## converged
## # weights: 9
## initial value 19001.093244
## final value 16026.916900
## converged
## # weights: 25
## initial value 18234.740751
## final value 16026.916900
```

```
## converged
## # weights: 41
## initial value 17444.773107
## final value 16026.916900
## converged
## # weights: 9
## initial value 19226.609443
## iter 10 value 16031.661108
## final value 16031.376945
## converged
## # weights: 25
## initial value 17912.050438
## iter 10 value 16099.176152
## iter 20 value 16031.084078
## iter 30 value 16030.127163
## final value 16030.125564
## converged
## # weights: 41
## initial value 19636.734646
## iter 10 value 16041.491939
## iter 20 value 16030.108778
## iter 30 value 16028.945067
## iter 40 value 16028.760813
## iter 50 value 16028.737686
## final value 16028.735586
## converged
## # weights: 9
## initial value 17894.055196
## iter 10 value 16037.383250
## iter 20 value 16027.037569
## iter 20 value 16027.037521
## final value 16027.037521
## converged
## # weights: 25
## initial value 17695.042357
## iter 10 value 16036.926838
## iter 20 value 16027.032307
## final value 16026.947386
## converged
## # weights: 41
## initial value 16889.040753
## iter 10 value 16033.782658
## iter 20 value 16026.996057
## final value 16026.924364
## converged
## # weights: 9
## initial value 18092.146826
## final value 16199.115500
## converged
## # weights: 25
```

```
## initial value 18173.879067
## final value 16199.115500
## converged
## # weights: 41
## initial value 19348.973466
## final value 16199.115500
## converged
## # weights: 9
## initial value 17686.176332
## iter 10 value 16205.168079
## final value 16203.578217
## converged
## # weights: 25
## initial value 18221.044579
## iter 10 value 16206.649650
## iter 20 value 16201.780743
## iter 30 value 16201.659845
## final value 16201.653530
## converged
## # weights: 41
## initial value 18007.644613
## iter 10 value 16236.916646
## iter 20 value 16201.368158
## iter 30 value 16200.961996
## final value 16200.935019
## converged
## # weights: 9
## initial value 18647.736276
## iter 10 value 16206.810589
## iter 20 value 16199.204218
## iter 20 value 16199.204183
## final value 16199.204183
## converged
## # weights: 25
## initial value 17681.544301
## iter 10 value 16203.519924
## iter 20 value 16199.166280
## final value 16199.138361
## converged
## # weights: 41
## initial value 18978.786685
## iter 10 value 16214.630998
## iter 20 value 16199.294382
## iter 30 value 16199.129205
## iter 30 value 16199.129191
## final value 16199.128531
## converged
## # weights: 9
## initial value 18177.074240
## final value 15821.871200
```

```
## converged
## # weights: 25
## initial value 18778.113264
## final value 15821.871200
## converged
## # weights: 41
## initial value 16909.381852
## final value 15821.871200
## converged
## # weights: 9
## initial value 17119.260930
## iter 10 value 15826.599396
## iter 20 value 15826.324908
## final value 15826.324509
## converged
## # weights: 25
## initial value 18703.665234
## iter 10 value 15832.759440
## iter 20 value 15826.738148
## iter 30 value 15824.517679
## iter 40 value 15824.417539
## iter 50 value 15824.405208
## final value 15824.404153
## converged
## # weights: 41
## initial value 18070.771384
## iter 10 value 15824.511212
## iter 20 value 15823.959776
## iter 30 value 15823.687857
## final value 15823.686523
## converged
## # weights: 9
## initial value 18880.142624
## iter 10 value 15828.619435
## iter 20 value 15821.949002
## iter 20 value 15821.948971
## final value 15821.948971
## converged
## # weights: 25
## initial value 17263.346538
## iter 10 value 15823.989488
## iter 20 value 15821.994572
## final value 15821.983647
## converged
## # weights: 41
## initial value 17737.097268
## iter 10 value 15841.716075
## iter 20 value 15822.099996
## final value 15821.883136
## converged
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info =  
trainInfo,  
## : There were missing values in resampled performance measures.
```

```
## # weights: 9  
## initial value 23385.579068  
## iter 10 value 20079.130735  
## iter 20 value 20063.354114  
## iter 20 value 20063.354040  
## final value 20063.354040  
## converged
```

NeuralNet_Cars

```
## Neural Network  
##  
## 400 samples  
## 6 predictor  
##  
## No pre-processing  
## Resampling: Cross-Validated (5 fold, repeated 2 times)  
## Summary of sample sizes: 320, 321, 319, 320, 320, 319, ...  
## Resampling results across tuning parameters:
```

```
##  
## size decay RMSE Rsquared MAE  
## 1 0e+00 7.080831 NaN 6.511777  
## 1 1e-04 7.080831 NaN 6.511777  
## 1 1e-01 7.081018 0.02150646 6.511976  
## 3 0e+00 7.080831 NaN 6.511777  
## 3 1e-04 7.080831 0.05189222 6.511777  
## 3 1e-01 7.080934 0.03272925 6.511887  
## 5 0e+00 7.080831 NaN 6.511777  
## 5 1e-04 7.080831 0.03626897 6.511777  
## 5 1e-01 7.080900 0.02608450 6.511851  
##
```

```
## RMSE was used to select the optimal model using the smallest value.  
## The final values used for the model were size = 1 and decay = 1e-04.
```

#The optimal model was chosen based on the RMSE (root mean squared error), featuring a size of 1 and a decay of 0. The Root mean squared error for this model was 7.081637, combined with a corresponding MAE (mean absolute error) value of 6.511536. While the Rsquared value was 'not available' for the optimal model, it ranged from NaN to 0.02538470 for other models.

#QB4 - "Consider the following input: Sales=9, Price=6.54, Population=124, Advertising=0, Age=76, Income= 110, Education=10 What will be the estimated Sales for this record using the above neuralnet model?"

```
Sales <- c(9)  
Price <- c(6.54)  
Population <- c(124)
```



```
Advertising <- c(0)
Age <- c(76)
Income <- c(110)
Education <- c(10)
Neural_Net_Testing <- data.frame(Sales, Price, Population, Advertising, Age,
Income, Education)

Sales_Prediction <- predict(NeuralNet_Cars, newdata = Neural_Net_Testing)
Sales_Prediction

## 1
## 1
```